

## Exercises – Third week

### 1. $\epsilon$ -greedy method on the 10-armed bandit problem

#### **Task:**

Implement the  $\epsilon$ -greedy algorithm for solving a 10-armed bandit problem with the following setup:

- $n = 10$  possible actions
- Each  $Q^*(a)$  is chosen randomly from a normal distribution:  $\eta(0; 1)$
- Each  $r_t$  is also normal:  $\eta(Q^*(a_t); 1)$
- 1000 plays
- Repeat the whole thing 2000 times and average the results

Run experiments with  $\epsilon = 0.1$ ;  $\epsilon = 0.01$  and  $\epsilon = 0.001$ . Finally, plot the average curves for each value of  $\epsilon$ .

#### **Implementation:**

```
clc
clear all
close all
action = zeros(10,1);
%rt = zeros(1000,1);
Q = [1:10];
Qstar = [1:10];
count = [1:10];
rmean = [1:1001];
rmean(1) = 0;

for i = 1:1:10
    Qstar(i) = normrnd(0,1);
end
epsvec = [1:3];
epsvec(1) = 0.1;
epsvec(2) = 0.01;
epsvec(3) = 0.001;

rmeanmat = zeros(3,1001);

for k=1:1:3
    eps = epsvec(k);
    for i=1:1:1000
        rmean(i+1) = 0;
    end
    for m = 1:1:2000
        for n = 1:1:10
            Q(n) = 0;
            count(n) = 0;
        end

        for j = 1:1:1000;

            prob = rand();
            if prob < (1-eps)
                [~,a] = max(Q);
            else
                a = randi([1,10],1,1);
            end
            count(a) = count(a) + 1;
            r = normrnd(Qstar(a),1);
            rmean(j+1) = rmean(j+1) + r;
            Q(a) = Q(a)+ 1/count(a)*(r-Q(a));

        end
    end
end
```

```
end
for i=1:1:1001
    rmeanmat(k,i) = rmean(i)/2000;
end
end

x=[0:1000];
figure
plot(x,rmeanmat(1,:), 'c', x,rmeanmat(2,:), 'b', x,rmeanmat(3,:), 'r');
```

### **Results:**




Figure 1: Average depicts curves for  $\epsilon = 0.1$ ;  $\epsilon = 0.01$  and  $\epsilon = 0.009$ .

## 2. Iterative Policy Evaluation

### **Tasks:**

- 1) Implement the Iterative Policy Evaluation.
- 2) Perform experiments with the problem of A Small Gridworld depicted in page 5 of the presentation on dynamic programming.
- 3) In the report for this exercise you should include the Final Value Function.

### **Implementation:**

```
function [V] = iter_poly_eval()

gamma = 1;
sideL = 4;
nGrids = sideL^2;
V = zeros(sideL);
MAX_N_ITERS = 100; iterCnt = 0;
CONV_TOL = 1e-6; delta = 1e10;
pol_pi = 0.25;

while((delta > CONV_TOL) && (iterCnt <= MAX_N_ITERS))
    delta = 0;

    for ii=1:sideL,
        for jj=1:sideL,
            if((ii==1 && jj==1) || (ii==sideL && jj==sideL)) continue; end
```

```

v      = V(ii,jj);
v_tmp = 0.0;

% action = UP
if(ii==1)
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj));
elseif(ii==2 && jj==1)
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii-1,jj));
else
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii-1,jj));
end

% action = DOWN
if( ii==sideL )
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj));
elseif( ii==sideL-1 && jj==sideL )
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii+1,jj));
else
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii+1,jj));
end

% action = RIGHT
if( jj==sideL )
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj));
elseif( jj==sideL-1 && ii==sideL )
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj+1));
else
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj+1));
end

% action = LEFT
if( jj==1 )
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj));
elseif( jj==2 && ii==1 )
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj-1));
else
    v_tmp = v_tmp + pol_pi*(-1 + gamma*V(ii,jj-1));
end

% update V(ii,jj):
V(ii,jj) = v_tmp;

    delta = max([delta, abs(v-V(ii,jj))]);
end
end

iterCnt=iterCnt+1;
if(0 && mod(iterCnt,1)==0)
    fprintf('iterCnt=%5d; delta=%10.5f\n', iterCnt, delta);
    disp(round(V*10)/10);
    pause
end
end
end

```

### **Result:**

Final Value Function

0.00000	-13.99785	-19.99692	-21.99661
-13.99785	-17.99737	-19.99714	-19.99717
-19.99692	-19.99714	-17.99759	-13.99820
-21.99661	-19.99717	-13.99820	0.00000